

Epkserimenti

Friday, July 6, 2018 10:41 AM

Podaci su podeljeni u tri grupe:

- Train ~28000
- Validation ~4000
- Test ~4000

Data:

- 0: -4593 images- *Angry*
- 1: -547 images- *Disgust*
- 2: -5121 images- *Fear*
- 3: -8989 images- *Happy*
- 4: -6077 images- *Sad*
- 5: -4002 images- *Surprise*
- 6: -6198 images- *Neutral*

Primerak podataka:





1.
Conv2D(64, 3, relu, l2(0.01))

2x Conv2D(128, 3)
BatchNorm()
Relu()
MaxPool(2, 2)

2x Conv2D(512, 3, relu)
BatchNorm()
Relu()
MaxPool(2, 2)
Dropout(0.5)

Dense(512)
Dropout(0.5)
Relu()

Dense(7, softmax)



2.
Conv2D(64, 3, relu)



2x Conv2D(128, 3)
BatchNorm()
Relu()
MaxPool(2, 2)

Conv2D(512, 3, relu)
BatchNorm()
Relu()
MaxPool(2, 2)

Conv2D(512, 3, relu, l2(0.001))
BatchNorm()
Relu()
MaxPool(2, 2)

Dense(512)
Relu()

Dense(7, softmax)





Adagrad
Adadelata

SGD Default
loss: 1.3323 - acc: 0.6309 - val_loss: 1.4286 - val_acc: 0.6016

Adadelata isto samo malo brze


Adagrad zasad sporije, ali mozda ima nade da ode dalje, psticemo jos 100

```
model = Sequential()
model.add(Conv2D(
    input_shape=(48, 48, 1),
    filters=64,
    kernel_size=(3, 3),
    kernel_initializer=VarianceScaling(),
    activation='relu',
    trainable=False
))
model.add(MaxPooling2D(
    pool_size=(2, 2),
    strides=(2, 2)
))
model.add(Conv2D(
    filters=128,
    kernel_size=(3, 3),
    kernel_initializer=VarianceScaling(),
    use_bias=False,
    trainable=False
))
model.add(BatchNormalization(axis=-1))
model.add(Activation('relu'))
model.add(MaxPooling2D(
    pool_size=(2, 2),
    strides=(2, 2)
))
model.add(Conv2D(
    filters=512,
    kernel_size=(3, 3),
    kernel_initializer=VarianceScaling(),
    use_bias=False,
    trainable=False
))
model.add(BatchNormalization(axis=-1))
model.add(Activation('relu'))
model.add(MaxPooling2D(
    pool_size=(2, 2),
    strides=(2, 2)
))
model.add(Conv2D(
    filters=512,
    kernel_size=(3, 3),
    kernel_initializer=VarianceScaling(),
    kernel_regularizer=regularizers.l2(0.001),
    use_bias=False,
    trainable=False
))
model.add(BatchNormalization(axis=-1))
model.add(Activation('relu'))
model.add(MaxPooling2D(
    pool_size=(2, 2),
    strides=(2, 2)
))
model.add(Flatten())
model.add(Dense(
    units=512,
    kernel_initializer=VarianceScaling(),
    activation='relu'
))
model.add(Dense(
    units=256,
    kernel_initializer=VarianceScaling(),
    activation='relu'
))
model.add(Dense(
    units=128,
    kernel_initializer=VarianceScaling(),
    activation='relu'
))
model.add(Dense(
    units=64,
    kernel_initializer=VarianceScaling(),
    activation='relu'
))
model.add(Dense(
    units=32,
    kernel_initializer=VarianceScaling(),
    activation='relu'
))
model.add(Dense(
    units=16,
    kernel_initializer=VarianceScaling(),
    activation='relu'
))
model.add(Dense(
    units=7,
    kernel_initializer=VarianceScaling(),
    activation='softmax'
))
model.compile(
    loss='categorical_crossentropy',
    optimizer="adam",
    metrics=[ 'accuracy' ]
)
```



Pokusacemo da zakljucamo conv layer I da nastavimo trening nad denseovima
U prvih par iteracija odmah se malo poboljsao val_acc

```
Epoch 1/200
50/50 [=====] - 6s 122ms/step - loss: 0.8023 - acc: 0.8120 - val_loss: 1.4885 - val_acc: 0.6403
Epoch 2/200
50/50 [=====] - 5s 99ms/step - loss: 0.8134 - acc: 0.8069 - val_loss: 1.4264 - val_acc: 0.6451
Epoch 3/200
50/50 [=====] - 5s 97ms/step - loss: 0.7651 - acc: 0.8225 - val_loss: 1.5053 - val_acc: 0.6471
Epoch 4/200
50/50 [=====] - 5s 95ms/step - loss: 0.7950 - acc: 0.8158 - val_loss: 1.5222 - val_acc: 0.6363
Epoch 5/200
50/50 [=====] - 5s 104ms/step - loss: 0.7798 - acc: 0.8191 - val_loss: 1.4964 - val_acc: 0.6419
Epoch 6/200
50/50 [=====] - 5s 97ms/step - loss: 0.7634 - acc: 0.8217 - val_loss: 1.5499 - val_acc: 0.6440
Epoch 7/200
50/50 [=====] - 5s 97ms/step - loss: 0.7661 - acc: 0.8223 - val_loss: 1.5077 - val_acc: 0.6410
Epoch 8/200
50/50 [=====] - 5s 97ms/step - loss: 0.7673 - acc: 0.8248 - val_loss: 1.5227 - val_acc: 0.6420
Epoch 9/200
50/50 [=====] - 5s 100ms/step - loss: 0.7687 - acc: 0.8223 - val_loss: 1.4808 - val_acc: 0.6394
Epoch 10/200
50/50 [=====] - 5s 94ms/step - loss: 0.7572 - acc: 0.8255 - val_loss: 1.5059 - val_acc: 0.6527
Epoch 11/200
50/50 [=====] - 5s 98ms/step - loss: 0.7668 - acc: 0.8256 - val_loss: 1.4796 - val_acc: 0.6423
Epoch 12/200
50/50 [=====] - 5s 100ms/step - loss: 0.7559 - acc: 0.8240 - val_loss: 1.4777 - val_acc: 0.6461
Epoch 13/200
```



Pokusacemo da zakljucamo prvi dense layer

Nema promene

Pokusacemo da odkljucamo konv layer da vidimo sta ce da se desi


Val_acc pao sa 0.65 na 0.6

)


Isti model kao iznad, povecan broj filtera u drugoj konvoluciji

```
model = Sequential()
model.add(Conv2D(
    input_shape=(48, 48, 1),
    filters=64,
    kernel_size=(3, 3),
    kernel_initializer=VarianceScaling(),
    activation='relu'
))
model.add(MaxPooling2D(
    pool_size=(2, 2),
    strides=(2, 2)
))
model.add(Conv2D(
    filters=256, <----- 128->256
    kernel_size=(3, 3),
    kernel_initializer=VarianceScaling(),
    use_bias=False
))
model.add(BatchNormalization(axis=-1))
model.add(Activation('relu'))
model.add(MaxPooling2D(
    pool_size=(2, 2),
    strides=(2, 2)
))
model.add(Conv2D(
    filters=512,
    kernel_size=(3, 3),
    kernel_initializer=VarianceScaling(),
    use_bias=False
))
model.add(BatchNormalization(axis=-1))
model.add(Activation('relu'))
model.add(MaxPooling2D(
    pool_size=(2, 2),
    strides=(2, 2)
))
model.add(Conv2D(
    filters=512,
    kernel_size=(3, 3),
    kernel_initializer=VarianceScaling(),
    kernel_regularizer=regularizers.l2(0.001),
    use_bias=False
))
model.add(BatchNormalization(axis=-1))
model.add(Activation('relu'))
model.add(MaxPooling2D(
    pool_size=(2, 2),
    strides=(2, 2)
))
model.add(Flatten())
model.add(Dense(
    units=512,
    kernel_initializer=VarianceScaling(),
    activation='relu'
))
model.add(Dense(
    units=256,
    kernel_initializer=VarianceScaling(),
    activation='relu'
))
model.add(Dense(
    units=128,
    kernel_initializer=VarianceScaling(),
    activation='relu'
))
model.add(Dense(
    units=64,
    kernel_initializer=VarianceScaling(),
    activation='relu'
))
model.add(Dense(
    units=32,
    kernel_initializer=VarianceScaling(),
    activation='relu'
))
model.add(Dense(
    units=16,
    kernel_initializer=VarianceScaling(),
    activation='relu'
))
model.add(Dense(
    units=7,
    kernel_initializer=VarianceScaling(),
    activation='softmax'
))
model.compile(
    loss='categorical_crossentropy',
    optimizer="adam",
    metrics=['accuracy']
)
```

acc




val_acc




Nakon zakljucavanja conv mreza

acc




val_acc




Dodali smo dropout na dense layer nakon treniranja, rezultat ovoga je bio sledeci (validation accuracy je konzistentniji)

acc




val_acc




Nakon ovoga, pokusali smo da odkljucamo konvolucione mreze, ali je val_acc pocao da opada (predpostavljamo da je do overfittinga)

acc




val_acc



Kad smo videli da val_acc opada, prekinuli smo treniranje da bi probali nesto drugo. Ovog puta stavljamo dropout na konv mreze i odkljucavamo ih.

acc



val_acc

